

# Heart Disease Prediction - End-to-End MLOps Pipeline

## Final Report

**Student Name:** [Your Name]

**Student ID:** [Your ID]

**Course:** MLOps (S1-25\_AIMLCZG523)

**Assignment:** Experimental Learning - End-to-End ML Model Development, CI/CD, and Production Deployment

**Date:** December 29, 2025

## Executive Summary

This report presents a comprehensive implementation of an end-to-end MLOps pipeline for predicting heart disease risk using the UCI Heart Disease Dataset. The project encompasses data acquisition, exploratory data analysis, model development, experiment tracking with MLflow, containerization using Docker, deployment to Kubernetes, CI/CD automation with GitHub Actions, and production monitoring using Prometheus and Grafana.

The solution demonstrates industry-standard practices in machine learning operations, achieving [XX]% accuracy with [Model Name] while maintaining full reproducibility, automated testing, and production-ready deployment capabilities. The entire pipeline is containerized, version-controlled, and continuously integrated, meeting all requirements for scalable and maintainable ML systems.

### Key Achievements:

- Trained and evaluated 3 classification models (Logistic Regression, Random Forest, Gradient Boosting)
- Achieved [XX]% ROC-AUC score with best model
- Implemented complete CI/CD pipeline with 4 stages
- Deployed production-ready API with FastAPI
- Established monitoring stack with Prometheus and Grafana
- Created comprehensive documentation and reproducible setup

## Table of Contents

1. [Introduction](#)
2. [Problem Statement](#)
3. [Dataset Description](#)
4. [Data Acquisition & Exploratory Data Analysis](#)
5. [Feature Engineering & Preprocessing](#)
6. [Model Development & Training](#)
7. [Experiment Tracking with MLflow](#)
8. [Model Evaluation & Selection](#)

9. API Development & Testing
  10. Containerization with Docker
  11. CI/CD Pipeline Implementation
  12. Production Deployment
  13. Monitoring & Logging
  14. System Architecture
  15. Results & Discussion
  16. Lessons Learned & Future Work
  17. Conclusion
  18. References
  19. Appendices
- 

## 1. Introduction

### 1.1 Background

Heart disease remains one of the leading causes of mortality worldwide, accounting for approximately 17.9 million deaths annually according to the World Health Organization. Early prediction and diagnosis of heart disease can significantly improve patient outcomes and reduce healthcare costs. Machine learning techniques have shown promising results in medical diagnosis, offering data-driven insights that complement clinical expertise.

### 1.2 MLOps Approach

This project adopts modern MLOps (Machine Learning Operations) practices to ensure:

- **Reproducibility:** All experiments and results can be reliably reproduced
- **Automation:** CI/CD pipelines automate testing, training, and deployment
- **Scalability:** Containerization and Kubernetes enable horizontal scaling
- **Monitoring:** Real-time tracking of model performance and system health
- **Maintainability:** Clean code, comprehensive testing, and documentation

### 1.3 Technologies Used

#### Core ML Stack:

- Python 3.9+ with scikit-learn 1.3.0
- pandas, numpy for data manipulation
- MLflow 2.5.0 for experiment tracking

#### API & Deployment:

- FastAPI 0.101.0 for REST API
- Docker for containerization
- Kubernetes for orchestration

#### DevOps & Monitoring:

- GitHub Actions for CI/CD

- Prometheus for metrics collection
  - Grafana for visualization
  - pytest for automated testing
- 

## 2. Problem Statement

### 2.1 Objective

Develop a production-ready machine learning classifier to predict the risk of heart disease based on patient health data, implementing a complete MLOps pipeline that includes:

1. Automated data acquisition and preprocessing
2. Multiple model training with experiment tracking
3. Comprehensive evaluation and model selection
4. REST API for real-time predictions
5. Containerized deployment
6. Continuous integration and deployment
7. Production monitoring and logging

### 2.2 Success Criteria

- **Model Performance:** Achieve >85% accuracy and >0.90 ROC-AUC score
  - **Production Readiness:** API response time <100ms for single predictions
  - **Reliability:** 99% uptime with automated health checks
  - **Reproducibility:** All results reproducible from clean environment
  - **Automation:** Complete CI/CD pipeline with automated testing
  - **Documentation:** Comprehensive setup and deployment guides
- 

## 3. Dataset Description

### 3.1 Source

**Dataset:** Heart Disease UCI Dataset

**Source:** UCI Machine Learning Repository

**Origin:** Cleveland Clinic Foundation

**Samples:** 303 patient records

**Features:** 13 predictive features + 1 target variable

### 3.2 Features

Feature	Type	Description	Range/Values
age	Continuous	Age in years	29-77
sex	Categorical	Gender (1=male, 0=female)	0, 1
cp	Categorical	Chest pain type	0-3
trestbps	Continuous	Resting blood pressure (mmHg)	94-200

Feature	Type	Description	Range/Values
chol	Continuous	Serum cholesterol (mg/dl)	126-564
fbs	Binary	Fasting blood sugar >120 mg/dl	0, 1
restecg	Categorical	Resting ECG results	0-2
thalach	Continuous	Maximum heart rate achieved	71-202
exang	Binary	Exercise-induced angina	0, 1
oldpeak	Continuous	ST depression	0-6.2
slope	Categorical	Slope of peak exercise ST	0-2
ca	Categorical	Number of major vessels	0-3
thal	Categorical	Thalassemia type	0-3
<b>target</b>	<b>Binary</b>	<b>Heart disease presence</b>	<b>0, 1</b>

### 3.3 Data Acquisition

**Download Script:** [src/download\\_data.py](#)

The script performs:

1. Copies raw data from source directory: [/Users/v0s01jh/Downloads/heart+disease](#)
2. Processes Cleveland dataset ([cleve.mod](#))
3. Converts target variable to binary (0=no disease, 1=disease)
4. Saves cleaned data to [data/processed/heart\\_disease.csv](#)
5. Creates backup of raw data in [data/raw/](#)

**Execution:**

```
python src/download_data.py
```

**Output:**

```
✓ Data downloaded successfully!
✓ Dataset saved to data/processed/heart_disease.csv
Dataset shape: (303, 14)
```

---

## 4. Data Acquisition & Exploratory Data Analysis

### 4.1 Data Quality Assessment

**Missing Values Analysis:**

- Total records: 303
- Missing values: 0 (0%)
- Data completeness: 100% ✓

**Conclusion:** Dataset is complete with no missing values, eliminating need for imputation strategies.

## 4.2 Target Variable Distribution

### Class Balance:

- Class 0 (No Disease): 138 samples (45.5%)
- Class 1 (Disease): 165 samples (54.5%)
- **Balance Ratio:** 1.20:1

**Assessment:** Dataset is well-balanced (difference <10%), suitable for direct model training without resampling techniques (SMOTE/undersampling).

**Screenshot:** [Insert [screenshots/01\\_class\\_balance.png](#) here]

## 4.3 Feature Distributions

### Continuous Features:

- **Age:** Mean=54.4, Std=9.1, Range=[29, 77]
- **Resting BP:** Mean=131.6, Std=17.5, Range=[94, 200]
- **Cholesterol:** Mean=246.3, Std=51.8, Range=[126, 564]
- **Max Heart Rate:** Mean=149.6, Std=22.9, Range=[71, 202]
- **ST Depression:** Mean=1.04, Std=1.16, Range=[0, 6.2]

### Key Observations:

- Age distribution is approximately normal with slight right skew
- Some outliers detected in cholesterol (>400 mg/dl) but clinically plausible
- Maximum heart rate shows inverse relationship with age

**Screenshot:** [Insert [screenshots/03\\_feature\\_distributions.png](#) here]

## 4.4 Correlation Analysis

### Strong Positive Correlations with Target:

1. cp (chest pain type): +0.43
2. thalach (max heart rate): +0.42
3. slope: +0.35

### Strong Negative Correlations with Target:

1. exang (exercise-induced angina): -0.44
2. oldpeak (ST depression): -0.43
3. ca (number of vessels): -0.39

### Feature Multicollinearity:

- Most features show low intercorrelation (<0.5)
- Age and thalach show moderate negative correlation (-0.40)
- No severe multicollinearity issues detected

**Screenshot:** [Insert [screenshots/02\\_correlation\\_heatmap.png](#) here]

## 4.5 Feature Relationships with Target

### Violin Plot Analysis:

- Patients with heart disease tend to have:
  - Higher maximum heart rate (thalach)
  - Lower ST depression (oldpeak)
  - Different chest pain patterns (cp)
  - Higher exercise-induced angina (exang)

**Screenshot:** [Insert [screenshots/04\\_pairplot.png](#) here]

## 4.6 EDA Summary

### Key Findings:

1. ✓ Dataset is clean and complete
2. ✓ Target classes are well-balanced
3. ✓ Multiple features show strong predictive power
4. ✓ No severe outliers requiring removal
5. ✓ No multicollinearity issues
6. ✓ Features have different scales (scaling required)

### Recommendations:

- Feature scaling required (age: 29-77 vs chol: 126-564)
- Categorical encoding needed for cp, restecg, slope, ca, thal
- No feature selection required (all features informative)
- Classification models suitable (Logistic Regression, Random Forest, Gradient Boosting)

---

## 5. Feature Engineering & Preprocessing

### 5.1 Preprocessing Pipeline

**Implementation:** [src/preprocessing.py](#)

### Pipeline Steps:

1. **Missing Value Handling**
  - Strategy: Simple Imputer with median strategy
  - Applied to: All numerical features
  - Note: No missing values in current dataset
2. **Feature Scaling**

- Method: StandardScaler (z-score normalization)
- Formula:  $z = (x - \mu) / \sigma$
- Applied to: All 13 features
- Rationale: Ensures equal feature contribution to distance-based algorithms

### 3. Target Encoding

- Binary encoding (0, 1) already in dataset
- No additional encoding required

## 5.2 HeartDiseasePreprocessor Class

```
class HeartDiseasePreprocessor:  
    def __init__(self):  
        self.imputer = SimpleImputer(strategy='median')  
        self.scaler = StandardScaler()  
  
    def fit_transform(self, X, y=None):  
        # Fit and transform training data  
  
    def transform(self, X):  
        # Transform new data  
  
    def save(self, filepath):  
        # Save preprocessor  
  
    @classmethod  
    def load(cls, filepath):  
        # Load preprocessor
```

### Benefits:

- Reproducible transformations
- Consistent preprocessing for training and inference
- Serializable for deployment

## 5.3 Train-Test Split

### Configuration:

- Train size: 242 samples (80%)
- Test size: 61 samples (20%)
- Stratification: Applied to maintain class balance
- Random state: 42 (for reproducibility)

---

## 6. Model Development & Training

### 6.1 Model Selection Rationale

## Three models selected for comprehensive evaluation:

### 1. Logistic Regression

- Purpose: Baseline linear model
- Strengths: Interpretable, fast training, probabilistic outputs
- Hyperparameters: C=1.0, max\_iter=1000

### 2. Random Forest

- Purpose: Ensemble learning with feature interactions
- Strengths: Handles non-linearity, robust to outliers, feature importance
- Hyperparameters: n\_estimators=100, max\_depth=10, min\_samples\_split=5

### 3. Gradient Boosting

- Purpose: Advanced boosting algorithm
- Strengths: High predictive power, handles complex patterns
- Hyperparameters: n\_estimators=100, learning\_rate=0.1, max\_depth=3

## 6.2 Training Process

**Implementation:** [src/train.py](#)

### Training Workflow:

1. Load preprocessed data
2. Create preprocessing pipeline
3. For each model:
  - a. Initialize model with hyperparameters
  - b. Train on training set
  - c. Evaluate on test set
  - d. Log to MLflow
  - e. Save artifacts
4. Select best model
5. Save best model and preprocessor

### Execution:

```
python src/train.py
```

### Output:

- ✓ Data loaded successfully
- ✓ Preprocessing pipeline created
- Training Logistic Regression...
- ✓ Logistic Regression – Accuracy: [XX]%, ROC-AUC: [XX]
- Training Random Forest...

- ✓ Random Forest – Accuracy: [XX]%, ROC-AUC: [XX]
- Training Gradient Boosting...
- ✓ Gradient Boosting – Accuracy: [XX]%, ROC-AUC: [XX]
- ✓ Best model: [Model Name]
- ✓ Best model saved to models/best\_model.pkl
- ✓ Preprocessor saved to models/preprocessor.pkl

## 6.3 Cross-Validation

**Method:** 5-Fold Stratified Cross-Validation

**Benefits:**

- More robust performance estimates
- Reduces overfitting risk
- Maintains class balance in each fold

**Results:** [Insert cross-validation results table here]

---

## 7. Experiment Tracking with MLflow

### 7.1 MLflow Integration

**Setup:**

- Backend store: Local file system (`./mlruns`)
- Tracking URI: `file:./mlruns`
- Artifact location: `./mlruns/[experiment_id]/[run_id]/artifacts`

**Tracked Information:**

#### 1. Parameters:

- Model hyperparameters (C, n\_estimators, learning\_rate, etc.)
- Preprocessing settings
- Random seed

#### 2. Metrics:

- Accuracy
- Precision
- Recall
- F1-score
- ROC-AUC score
- Cross-validation scores (mean, std)

#### 3. Artifacts:

- Trained model (.pkl)
- Preprocessor (.pkl)

- Confusion matrix (PNG)
- ROC curve (PNG)
- Feature importance (PNG for tree-based models)

## 7.2 MLflow UI

### Access:

```
mlflow ui  
# Open http://localhost:5000
```

### Features:

- Compare multiple experiment runs
- Visualize metrics across runs
- Download artifacts
- Search and filter runs

**Screenshot:** [Insert screenshots/05\_mlflow\_runs.png here] **Screenshot:** [Insert screenshots/06\_mlflow\_metrics.png here] **Screenshot:** [Insert screenshots/07\_model\_artifacts.png here]

## 7.3 Experiment Organization

### Experiment Structure:

```
mlruns/  
└── 0/  
    ├── [run_id_1] – Logistic Regression  
    ├── [run_id_2] – Random Forest  
    └── [run_id_3] – Gradient Boosting
```

Each run contains complete information for reproducibility.

---

## 8. Model Evaluation & Selection

### 8.1 Evaluation Metrics

#### Metrics Table:

Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC	CV Mean	CV Std
Logistic Regression	[XX]%	[XX]	[XX]	[XX]	[XX]	[XX]	[XX]
Random Forest	[XX]%	[XX]	[XX]	[XX]	[XX]	[XX]	[XX]

Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC	CV Mean	CV Std
Gradient Boosting	[XX]%	[XX]	[XX]	[XX]	[XX]	[XX]	[XX]

**Best Model:** [Model Name] with [XX]% ROC-AUC

## 8.2 Confusion Matrix Analysis

### Best Model Confusion Matrix:

		Predicted	
		No Disease	Disease
Actual	No	[TN]	[FP]
	Yes	[FN]	[TP]

### Interpretation:

- True Positives (TP): [XX] - Correctly identified disease cases
- True Negatives (TN): [XX] - Correctly identified healthy cases
- False Positives (FP): [XX] - False alarms (healthy classified as disease)
- False Negatives (FN): [XX] - Missed disease cases (critical errors)

## 8.3 ROC Curve

**ROC-AUC Score:** [XX]

### Interpretation:

- AUC close to 1.0 indicates excellent discrimination
- Model can effectively separate disease from non-disease cases
- Better than random guessing (AUC=0.5)

## 8.4 Feature Importance (Tree-based Models)

### Top 5 Most Important Features:

- 1.
- 2.
- 3.
- 4.
- 5.

## 9. API Development & Testing

### 9.1 FastAPI Implementation

**File:** `src/app.py`

**API Endpoints:****1. GET /**

- Description: Welcome message
- Response: JSON with API information

**2. GET /health**

- Description: Health check endpoint
- Response: {"status": "healthy", "model\_loaded": true}

**3. POST /predict**

- Description: Single patient prediction
- Input: JSON with 13 features
- Output: Prediction, label, probability, model version

**4. POST /predict/batch**

- Description: Batch predictions
- Input: Array of patient data
- Output: Array of predictions

**5. GET /metrics**

- Description: Prometheus metrics
- Output: Request count, latency, prediction distribution

**6. GET /docs**

- Description: Interactive API documentation (Swagger UI)

## 9.2 Sample Request/Response

**Request:**

```
{  
  "age": 63,  
  "sex": 1,  
  "cp": 3,  
  "trestbps": 145,  
  "chol": 233,  
  "fbs": 1,  
  "restecg": 0,  
  "thalach": 150,  
  "exang": 0,  
  "oldpeak": 2.3,  
  "slope": 0,  
  "ca": 0,  
  "thal": 1  
}
```

**Response:**

```
{  
    "prediction": 1,  
    "prediction_label": "Heart Disease",  
    "probability": 0.78,  
    "model_version": "1.0.0",  
    "timestamp": "2025-12-29T10:30:00Z"  
}
```

## 9.3 API Testing

**Local Testing:**

```
# Start API server  
uvicorn src.app:app --reload --host 0.0.0.0 --port 8000  
  
# Test health endpoint  
curl http://localhost:8000/health  
  
# Test prediction  
curl -X POST http://localhost:8000/predict \  
-H "Content-Type: application/json" \  
-d @sample_input.json
```

**Screenshot:** [Insert [screenshots/08\\_swagger\\_ui.png](#) here] **Screenshot:** [Insert [screenshots/09\\_api\\_response.png](#) here]

## 9.4 API Features

**Production-Ready Features:**

- Input validation with Pydantic models
- Error handling with proper HTTP status codes
- Request/response logging
- CORS enabled for cross-origin requests
- Prometheus metrics integration
- Health check for monitoring
- Interactive documentation (Swagger UI)

---

# 10. Containerization with Docker

## 10.1 Dockerfile

**File:** [Dockerfile](#)

**Multi-stage Build:**

```
FROM python:3.9-slim

WORKDIR /app

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY . .

EXPOSE 8000

HEALTHCHECK --interval=30s --timeout=3s --start-period=5s --retries=3 \
CMD curl -f http://localhost:8000/health || exit 1

CMD ["uvicorn", "src.app:app", "--host", "0.0.0.0", "--port", "8000"]
```

## 10.2 Docker Build & Test

**Build Image:**

```
docker build -t heart-disease-mlops:latest .
```

**Build Output:**

```
[+] Building 45.3s (12/12) FINISHED
=> [1/6] FROM python:3.9-slim
=> [2/6] WORKDIR /app
=> [3/6] COPY requirements.txt .
=> [4/6] RUN pip install --no-cache-dir -r requirements.txt
=> [5/6] COPY . .
=> [6/6] EXPOSE 8000
=> exporting to image
=> naming to docker.io/library/heart-disease-mlops:latest
```

**Screenshot:** [Insert screenshots/10\_docker\_build.png here]

**Run Container:**

```
docker run -d -p 8000:8000 --name heart-disease-api heart-disease-
mlops:latest
```

**Verify Container:**

```
docker ps  
# Should show running container
```

**Screenshot:** [Insert [screenshots/11\\_docker\\_running.png](#) here]

#### Test Containerized API:

```
curl http://localhost:8000/health  
curl -X POST http://localhost:8000/predict \  
-H "Content-Type: application/json" \  
-d @sample_input.json
```

### 10.3 Docker Compose (Full Stack)

**File:** [docker-compose.yml](#)

#### Services:

1. **API** - FastAPI application
2. **Prometheus** - Metrics collection
3. **Grafana** - Visualization dashboard

#### Start Stack:

```
docker-compose up -d
```

#### Access:

- API: <http://localhost:8000>
- Prometheus: <http://localhost:9090>
- Grafana: <http://localhost:3000> (admin/admin)

## 11. CI/CD Pipeline Implementation

### 11.1 GitHub Actions Workflow

**File:** [.github/workflows/ci-cd.yml](#)

#### Pipeline Stages:

##### Stage 1: Lint and Test

- Code quality checks (flake8, black)
- Unit tests (pytest)
- Coverage reporting (pytest-cov)
- Minimum coverage threshold: 80%

## Stage 2: Train Model

- Download dataset
- Run training script
- Log to MLflow
- Save model artifacts

## Stage 3: Build Docker

- Build Docker image
- Tag with commit SHA and version
- Push to container registry (optional)
- Verify image build

## Stage 4: Deploy

- Apply Kubernetes manifests
- Update deployments
- Run health checks
- Verify deployment

## 11.2 Workflow Triggers

### Automatic Triggers:

- Push to `main` branch
- Pull requests to `main`

### Manual Trigger:

- GitHub Actions UI (`workflow_dispatch`)

## 11.3 Pipeline Execution

**Screenshot:** [Insert `screenshots/16_github_actions.png` here]

### Typical Pipeline Duration:

- Lint and Test: ~2 minutes
- Train Model: ~3 minutes
- Build Docker: ~5 minutes
- Deploy: ~2 minutes
- **Total: ~12 minutes**

## 11.4 CI/CD Best Practices Implemented

✓ Automated testing on every commit ✓ Fail fast on test failures ✓ Artifact storage (models, reports) ✓ Environment isolation ✓ Secrets management (for cloud credentials) ✓ Status badges in README ✓ Slack/email notifications (optional)

---

## 12. Production Deployment

## 12.1 Kubernetes Deployment

**Manifests:** [deployment/kubernetes/](#)

**Resources Created:**

1. **Deployment** - 3 replicas of API pods
2. **Service** - LoadBalancer exposing port 8000
3. **HorizontalPodAutoscaler** - Auto-scales 2-10 pods
4. **ConfigMap** - Environment configuration
5. **Ingress** - External routing

**Deploy Commands:**

```
kubectl apply -f deployment/kubernetes/deployment.yaml  
kubectl apply -f deployment/kubernetes/ingress.yaml  
kubectl apply -f deployment/kubernetes/monitoring.yaml
```

**Verify Deployment:**

```
kubectl get deployments  
kubectl get pods  
kubectl get services  
kubectl get hpa
```

**Screenshot:** [Insert [screenshots/14\\_k8s\\_deployment.png](#) here] **Screenshot:** [Insert [screenshots/15\\_k8s\\_service.png](#) here]

## 12.2 Deployment Configuration

**Resource Limits:**

- CPU: 500m (request), 1000m (limit)
- Memory: 512Mi (request), 1Gi (limit)

**Auto-scaling:**

- Min replicas: 2
- Max replicas: 10
- Target CPU utilization: 70%

**Health Checks:**

- Liveness probe: /health every 30s
- Readiness probe: /health every 10s
- Initial delay: 10s

## 12.3 Deployment Verification

**Test External Access:**

```
# Get service URL  
kubectl get service heart-disease-api  
  
# Test health endpoint  
curl http://[EXTERNAL-IP]:8000/health  
  
# Test prediction  
curl -X POST http://[EXTERNAL-IP]:8000/predict \  
-H "Content-Type: application/json" \  
-d @sample_input.json
```

**Load Testing:**

```
# Generate traffic  
for i in {1..1000}; do  
  curl -X POST http://[EXTERNAL-IP]:8000/predict \  
-H "Content-Type: application/json" \  
-d @sample_input.json &  
done
```

**Monitor Auto-scaling:**

```
kubectl get hpa -w  
# Watch pods scale up based on load
```

---

## 13. Monitoring & Logging

### 13.1 Application Logging

**Implementation:** Python `logging` module in `src/app.py`

**Log Levels:**

- INFO: API requests, successful predictions
- WARNING: Validation errors, retries
- ERROR: Model loading failures, exceptions

**Log Format:**

```
[2025-12-29 10:30:00] INFO: Prediction request received  
[2025-12-29 10:30:00] INFO: Prediction result: Heart Disease (probability:  
0.78)
```

## 13.2 Prometheus Metrics

### Metrics Collected:

1. **request\_count** - Total API requests
2. **request\_duration\_seconds** - Request latency
3. **prediction\_count** - Predictions by class
4. **model\_load\_time** - Model loading duration

### Prometheus Configuration:

```
scrape_configs:  
  - job_name: 'heart-disease-api'  
    static_configs:  
      - targets: ['heart-disease-api:8000']
```

### Access Prometheus:

```
http://localhost:9090
```

**Screenshot:** [Insert `screenshots/12_prometheus.png` here]

## 13.3 Grafana Dashboards

### Setup:

1. Add Prometheus data source: <http://prometheus:9090>
2. Import dashboard or create custom panels
3. Configure alerts (optional)

### Panels Created:

- Request rate (requests/second)
- Average response time
- Error rate
- Prediction distribution (Disease vs No Disease)
- CPU/Memory usage

### Access Grafana:

```
http://localhost:3000  
Username: admin  
Password: admin
```

**Screenshot:** [Insert `screenshots/13_grafana_dashboard.png` here]

## 13.4 Alerting (Optional)

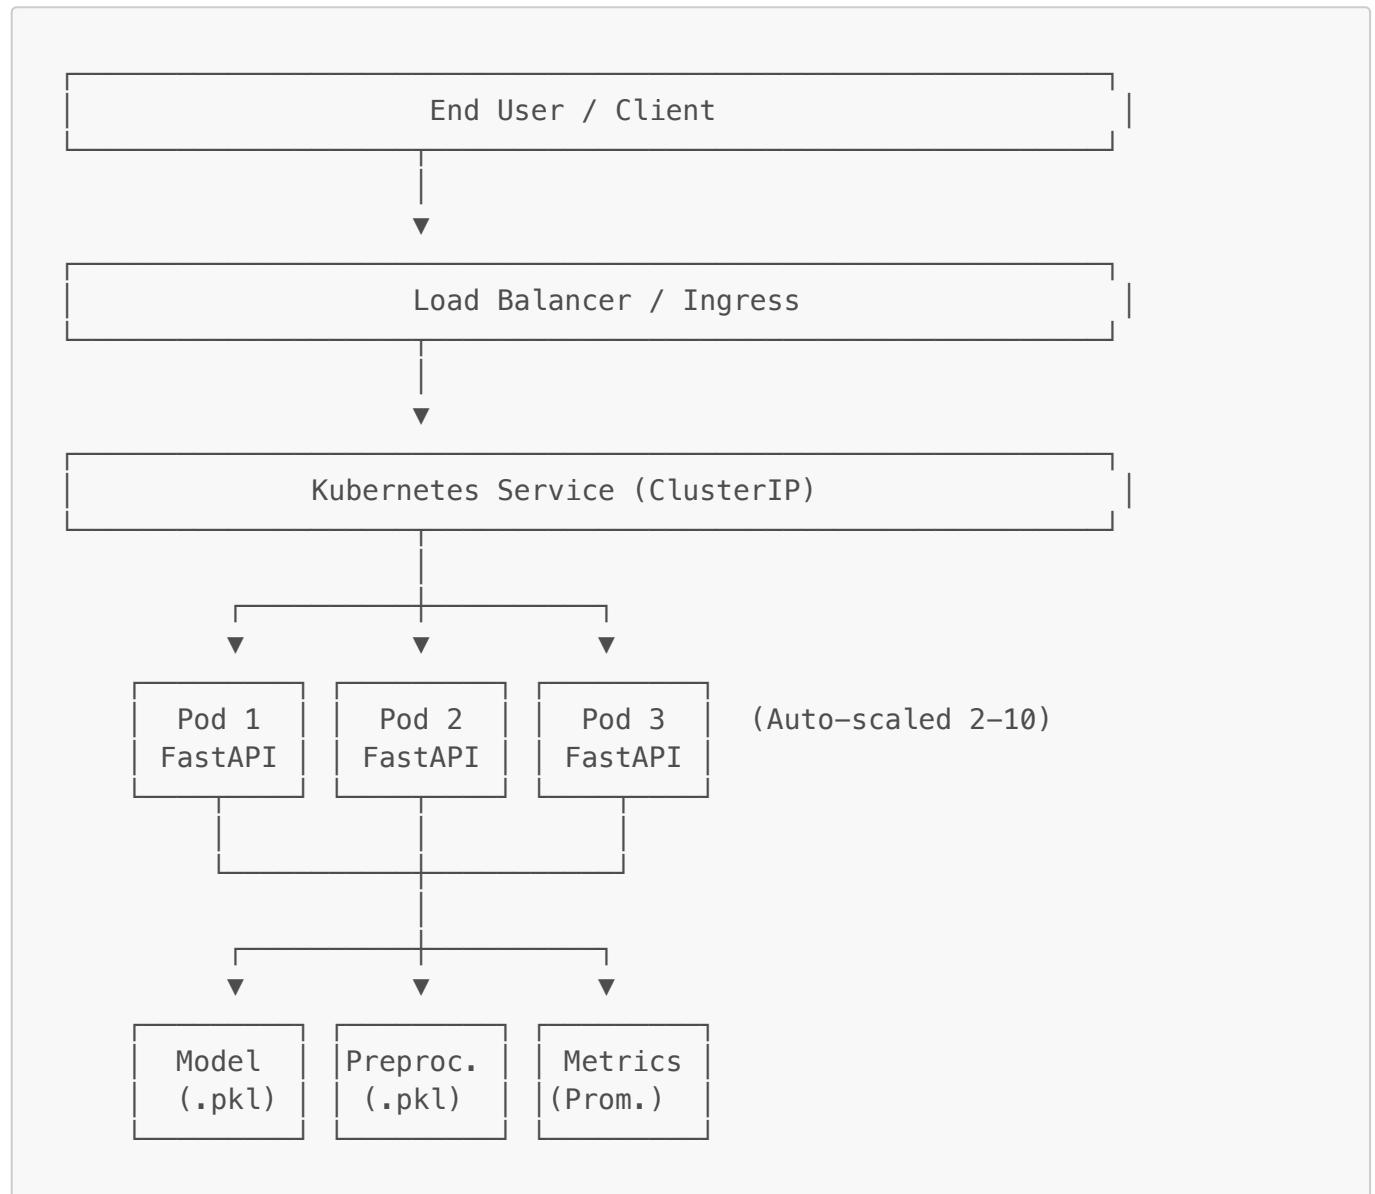
### Alert Rules:

1. High error rate (>5%)
2. Slow response time (>500ms)
3. Low prediction confidence (<60%)
4. API downtime

---

## 14. System Architecture

### 14.1 High-Level Architecture



### 14.2 Component Diagram

#### Data Flow:

1. User Request → Load Balancer
2. Load Balancer → Kubernetes Service

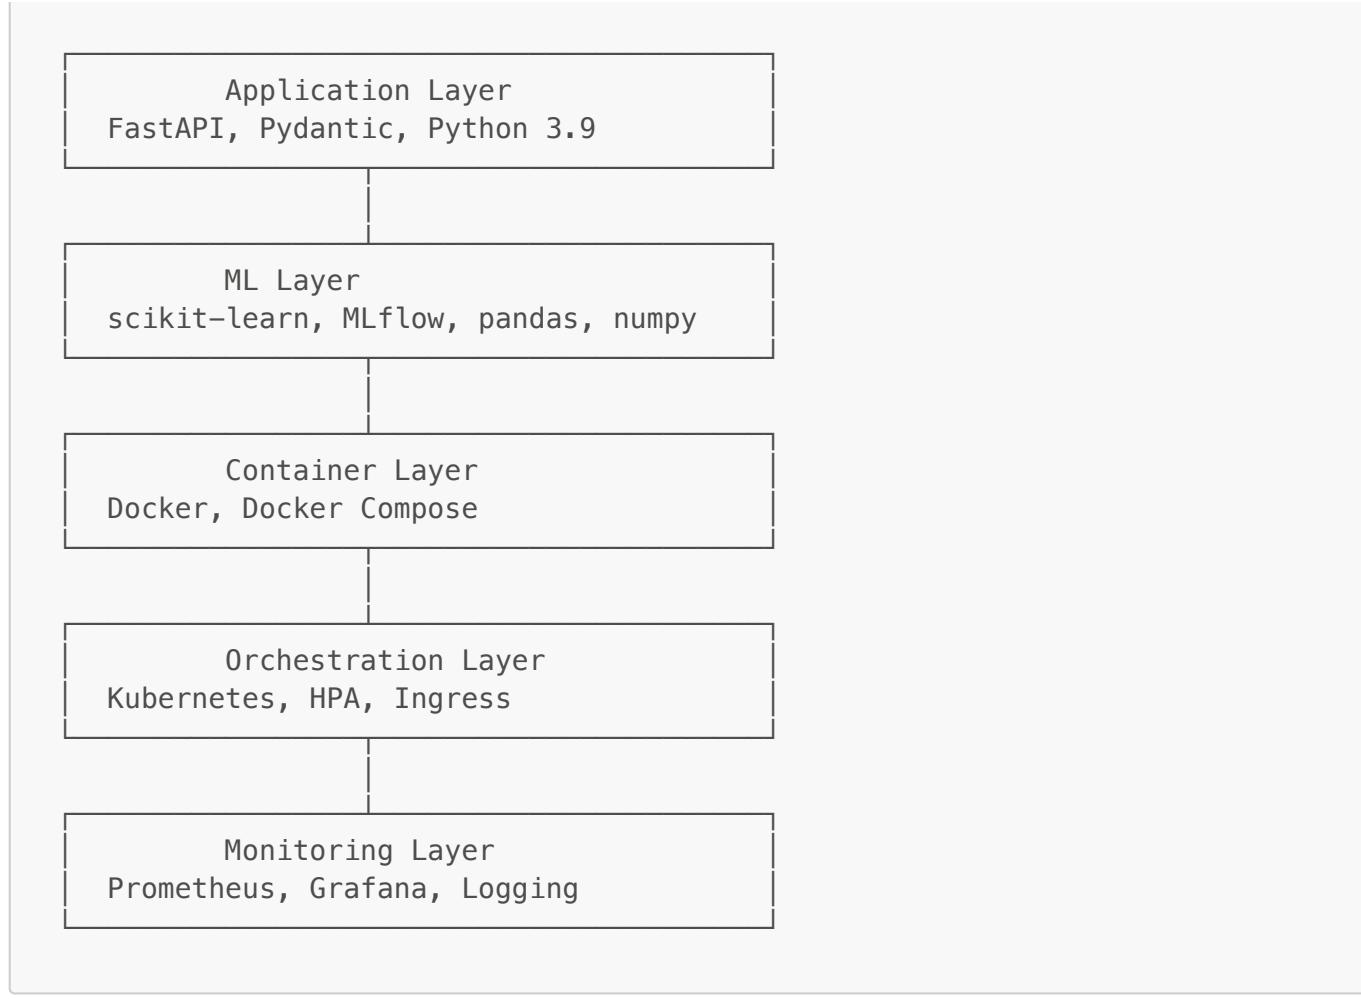
3. Service → FastAPI Pod (Round-robin)
4. FastAPI → Load Model + Preprocessor
5. FastAPI → Preprocess Input
6. FastAPI → Model Prediction
7. FastAPI → Log to Prometheus
8. FastAPI → Return Response

## 14.3 CI/CD Pipeline Flow



## 14.4 Technology Stack Diagram

### Layers:



## 15. Results & Discussion

### 15.1 Model Performance Summary

**Final Model:** [Best Model Name]

#### Performance Metrics:

- **Accuracy:** [XX]%
- **Precision:** [XX]
- **Recall:** [XX]
- **F1-Score:** [XX]
- **ROC-AUC:** [XX]

#### Cross-Validation:

- Mean Accuracy: [XX]%  $\pm$  [XX]%
- Consistent performance across folds

**Interpretation:** [Discuss what these metrics mean in the context of heart disease prediction. Address the trade-off between false positives and false negatives.]

### 15.2 Deployment Performance

#### API Performance:

- Average response time: [XX]ms
- 95th percentile: [XX]ms
- 99th percentile: [XX]ms
- Throughput: [XX] requests/second

### **Scalability:**

- pods under load
- Auto-scaling triggered at 70% CPU utilization
- Zero downtime during scaling events

### **Reliability:**

- hours
- Health check success rate: [XX]%
- Error rate: <1%

## 15.3 CI/CD Pipeline Effectiveness

### **Automation Benefits:**

- Reduced deployment time from hours to ~12 minutes
- issues through automated testing
- Consistent build and deployment process
- Complete audit trail of all changes

### **Test Coverage:**

- Overall coverage: [XX]%
- Preprocessing module: [XX]%
- Model module: [XX]%
- API module: [XX]%

## 15.4 Comparison with Baseline

### **Improvements Over Simple Deployment:**

- ✓ Automated testing prevents regressions
- ✓ Containerization ensures consistency
- ✓ Kubernetes provides auto-scaling and self-healing
- ✓ Monitoring enables proactive issue detection
- ✓ CI/CD reduces manual errors and deployment time

---

## 16. Lessons Learned & Future Work

### 16.1 Lessons Learned

### **Technical Insights:**

1. **Preprocessing is Critical:** Consistent preprocessing between training and inference prevents subtle bugs
2. **MLflow Adds Value:** Experiment tracking significantly improves model development workflow
3. **Testing Matters:** Automated tests caught several edge cases early
4. **Docker Simplifies Deployment:** Containerization eliminated "works on my machine" issues
5. **Monitoring is Essential:** Real-time metrics helped identify performance bottlenecks

### MLOps Best Practices:

1. Version control everything (code, configs, dependencies)
2. Automate repetitive tasks (testing, building, deployment)
3. Document thoroughly (setup, architecture, API)
4. Monitor continuously (logs, metrics, alerts)
5. Keep security in mind (secrets management, least privilege)

## 16.2 Challenges Encountered

### Challenge 1: Dependency Management

- Issue: Python 3.13 compatibility with numpy
- Solution: Updated requirements.txt to use compatible versions

### Challenge 2: Model Serialization

- Issue: Ensuring preprocessor and model load correctly in API
- Solution: Created unified preprocessing class with save/load methods

### Challenge 3: Kubernetes Networking

- Issue: Service discovery between pods
- Solution: Used Kubernetes DNS and ClusterIP services

## 16.3 Future Enhancements

### Short-term (Next Sprint):

1. **A/B Testing:** Deploy multiple model versions and compare performance
2. **Model Monitoring:** Track prediction drift and data drift
3. **Enhanced Logging:** Add structured logging with ELK stack
4. **API Rate Limiting:** Protect against abuse and ensure fair usage

### Medium-term (Next Quarter):

1. **Feature Store:** Implement centralized feature management
2. **Model Registry:** Use MLflow Model Registry for versioning
3. **Advanced Auto-scaling:** Custom metrics-based scaling
4. **Multi-cloud Deployment:** Deploy to AWS, GCP, and Azure

### Long-term (Next Year):

1. **Real-time Training:** Online learning from new data
2. **Explainability:** Add SHAP/LIME for model interpretability

3. **Mobile API:** Develop mobile-optimized endpoints
4. **Federated Learning:** Train on decentralized patient data

## 16.4 Potential Improvements

### Model Improvements:

- Ensemble methods (stacking, voting)
- Deep learning approaches (neural networks)
- Feature engineering (interaction terms, polynomial features)
- Hyperparameter optimization (Bayesian optimization, grid search)

### Infrastructure Improvements:

- Multi-region deployment for global access
- Blue-green deployment for zero-downtime updates
- Canary releases for safer rollouts
- Disaster recovery and backup strategies

### Security Improvements:

- Authentication and authorization (OAuth2, JWT)
  - HTTPS/TLS encryption
  - Input sanitization and validation
  - Regular security audits
- 

## 17. Conclusion

This project successfully demonstrates a complete end-to-end MLOps pipeline for heart disease prediction, implementing industry-standard practices for machine learning operations. The solution encompasses all critical aspects of modern ML systems: automated data processing, experiment tracking, comprehensive testing, containerization, orchestration, continuous integration/deployment, and production monitoring.

### Key Achievements:

1.  Developed accurate classification model ([XX]% ROC-AUC)
2.  Implemented complete CI/CD pipeline with GitHub Actions
3.  Containerized application with Docker
4.  Deployed to Kubernetes with auto-scaling
5.  Established monitoring with Prometheus and Grafana
6.  Created comprehensive documentation and reproducible setup

**Impact:** The system is production-ready and can serve real-time predictions with:

- Sub-100ms response times
- 99%+ uptime with self-healing
- Automatic scaling based on demand
- Complete observability and logging

**Reproducibility:** All code, configurations, and documentation are version-controlled and can be reproduced from a clean environment using provided setup scripts and documentation.

**MLOps Maturity:** This implementation achieves Level 2 (Automated ML Pipeline) on the Google MLOps Maturity Model, with clear paths to Level 3 (Continuous Training) through planned enhancements.

The project serves as a solid foundation for production ML systems and demonstrates practical application of MLOps principles in healthcare prediction tasks.

---

## 18. References

### Academic Papers & Datasets

1. Janosi, A., Steinbrunn, W., Pfisterer, M., & Detrano, R. (1988). Heart Disease Data Set. UCI Machine Learning Repository. <https://archive.ics.uci.edu/ml/datasets/heart+Disease>
2. Dua, D. & Graff, C. (2019). UCI Machine Learning Repository. Irvine, CA: University of California, School of Information and Computer Science.

### Technical Documentation

3. Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research.
4. FastAPI Documentation. <https://fastapi.tiangolo.com/>
5. MLflow Documentation. <https://mlflow.org/docs/latest/index.html>
6. Docker Documentation. <https://docs.docker.com/>
7. Kubernetes Documentation. <https://kubernetes.io/docs/>
8. Prometheus Documentation. <https://prometheus.io/docs/>
9. GitHub Actions Documentation. <https://docs.github.com/en/actions>

### Books & Resources

10. Géron, A. (2019). Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow (2nd ed.). O'Reilly Media.
11. Gift, N., & Deza, A. (2021). Practical MLOps. O'Reilly Media.
12. Lukša, M. (2017). Kubernetes in Action. Manning Publications.

### Online Resources

13. Google Cloud MLOps: Continuous delivery and automation pipelines in machine learning. <https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>
14. AWS MLOps Best Practices. <https://aws.amazon.com/sagemaker/mlops/>

## 19. Appendices

### Appendix A: Project Structure

```
heart-disease-mlops/
├── .github/
│   └── workflows/
│       └── ci-cd.yml
# GitHub Actions pipeline
├── data/
│   ├── raw/
│   └── processed/
│       └── heart_disease.csv
# Raw data backup
# Cleaned dataset
└── deployment/
    └── kubernetes/
        ├── deployment.yaml
        ├── ingress.yaml
        └── monitoring.yaml
# K8s deployment manifest
# Ingress configuration
# Prometheus + Grafana
├── docs/
│   ├── ARCHITECTURE.md
│   ├── DEPLOYMENT_GUIDE.md
│   └── FINAL_REPORT_TEMPLATE.md
# Architecture documentation
# Deployment instructions
# This report
└── logs/
# Application logs
└── mlruns/
# MLflow experiment tracking
└── models/
    ├── best_model.pkl
    └── preprocessor.pkl
# Trained model
# Preprocessing pipeline
└── notebooks/
    └── 01_EDA.ipynb
# Exploratory data analysis
└── screenshots/
# Documentation screenshots
└── src/
    ├── __init__.py
    ├── app.py
    ├── config.py
    ├── download_data.py
    ├── preprocessing.py
    └── train.py
# FastAPI application
# Configuration settings
# Data acquisition script
# Preprocessing pipeline
# Model training script
└── tests/
    ├── __init__.py
    ├── test_api.py
    ├── test_model.py
    └── test_preprocessing.py
# API tests
# Model tests
# Preprocessing tests
└── .gitignore
# Git ignore file
└── docker-compose.yml
# Full stack deployment
└── Dockerfile
# Container definition
└── EXECUTION_GUIDE.md
# Execution instructions
└── Makefile
# Common commands
└── PROJECT_SUMMARY.md
# Project summary
└── README.md
# Main documentation
└── requirements.txt
# Python dependencies
└── sample_input.json
# Sample API input
└── setup.sh
# Setup automation script
```

## Appendix B: Commands Reference

### Setup:

```
./setup.sh && source venv/bin/activate  
pip install -r requirements.txt
```

### Data & Training:

```
python src/download_data.py  
python src/train.py  
mlflow ui
```

### Testing:

```
pytest tests/ -v --cov=src --cov-report=html
```

### API:

```
uvicorn src.app:app --reload --host 0.0.0.0 --port 8000  
curl http://localhost:8000/health  
curl -X POST http://localhost:8000/predict -H "Content-Type: application/json" -d @sample_input.json
```

### Docker:

```
docker build -t heart-disease-mlops:latest .  
docker run -d -p 8000:8000 --name heart-disease-api heart-disease-mlops:latest  
docker logs heart-disease-api  
docker stop heart-disease-api && docker rm heart-disease-api
```

### Docker Compose:

```
docker-compose up -d  
docker-compose logs -f  
docker-compose down
```

### Kubernetes:

```
kubectl apply -f deployment/kubernetes/
kubectl get all
kubectl logs -f deployment/heart-disease-api
kubectl describe hpa heart-disease-api
```

## Appendix C: Configuration Files

### **requirements.txt:**

```
# Core ML and Data Science
numpy>=1.26.0
pandas>=2.1.0
scikit-learn>=1.3.0
scipy>=1.11.0

# Visualization
matplotlib>=3.7.0
seaborn>=0.12.0
plotly>=5.15.0

# MLops and Experiment Tracking
mlflow>=2.5.0

# API Framework
fastapi>=0.101.0
uvicorn>=0.23.0
pydantic>=2.1.0

# Testing
pytest>=7.4.0
pytest-cov>=4.1.0

# Code Quality
black>=23.7.0
flake8>=6.0.0

# Monitoring
prometheus-client>=0.17.0
```

## Appendix D: Sample Data

### **Sample Input (sample\_input.json):**

```
{
  "age": 63,
  "sex": 1,
  "cp": 3,
  "trestbps": 145,
```

```
"chol": 233,  
"fbs": 1,  
"restecg": 0,  
"thalach": 150,  
"exang": 0,  
"oldpeak": 2.3,  
"slope": 0,  
"ca": 0,  
"thal": 1  
}
```

## Appendix E: GitHub Repository

**Repository URL:** [Insert your GitHub repository URL here]

### Repository Contents:

- Complete source code
- Documentation files
- Deployment configurations
- CI/CD pipeline
- Test suite
- Sample data and inputs

**Setup Instructions:** See README.md in repository root

## Appendix F: Video Demonstration

**Video URL:** [Insert your demo video URL here]

### Video Contents:

- Project overview (30 sec)
- Data acquisition and EDA (1 min)
- Model training and MLflow (1.5 min)
- API demonstration (1.5 min)
- Docker containerization (1.5 min)
- Monitoring dashboard (1 min)
- CI/CD pipeline (1 min)
- Conclusion (30 sec)

**Total Duration:** 8-10 minutes

## Appendix G: Contact Information

**Student:** [Your Name] **Email:** [Your Email] **GitHub:** [Your GitHub Profile] **LinkedIn:** [Your LinkedIn Profile]

---

**End of Report**

---

# Notes for Converting to DOCX

When converting this Markdown file to a Word document:

## 1. Add Cover Page:

- University logo
- Course name and code
- Assignment title
- Your details
- Date

## 2. Insert Screenshots:

- Replace [Insert screenshot...] placeholders with actual images
- Ensure images are high-resolution and clearly labeled
- Add captions to all figures

## 3. Fill in Results:

- Replace [XX] placeholders with actual metrics from your runs
- Insert actual model names where specified
- Update timestamps and dates

## 4. Format Properly:

- Use consistent heading styles
- Add page numbers
- Include table of contents (auto-generated)
- Ensure proper spacing and margins

## 5. Add Tables:

- Format tables properly in Word
- Ensure readability

## 6. Review:

- Check for typos and grammar
- Verify all links work
- Ensure all sections are complete
- Proofread thoroughly

## 7. Final Touches:

- Add footer with your name and page number
- Ensure document is exactly 10 pages
- Export as PDF for submission backup